

OSH_LAB_01

赵瑞 PB16120853

环境: Ubuntu 16.04, QEMU 2.12.0-rc1, kernel Linux-4.15.14, busybox-1.28.1, GDB 8.1

调试跟踪工具安装

1. 下载[kernel](#)、[busybox](#)、[QEMU](#)

都是下载源码编译的。

这三个的版本匹配真的很重要，刚开始我装QEMU是直接 `apt install qemu`，装的QEMU版本比较旧，于是后面在GDB调试的时候就出现了问题

2. 安装QEMU，先解压，然后如下：

```
1 $ cd qemu-2.12.0-rc1/
2 $ sudo make clean
3 $ ./configure
4 $ sudo make
5 $ sudo make install
```

为了后面方便，把QEMU加入path

```
1 ln -s /usr/bin/qemu-system-x86_64 /usr/bin/qemu
```

3. 安装busybox，也是先解压，然后如下安装

注意！：在第二行命令时，要修改一些东西

因为Linux运行环境当中是不带动态库的，所以必须以静态方式来编译Busybox。修改

Busybox Settings ---> Build Options ---> [*] Build Busybox as a static binary(no shared libs)

```
1 $ sudo make defconfig
2 $ sudo make menuconfig
3 $ sudo make
4 $ sudo make install
```

4. 编译内核（注意：在 `make menuconfig` 时候要把debug的东西选上。

```

1 ~$ cd linux/
2 ~/linux$ ls
3 build-initrd.sh  busybox-1.28.2.tar.bz2  linux-4.15.14.tar.gz
4 busybox-1.28.2  linux-4.15.14
5 $ cd linux-4.15.14/
6 $ sudo su
7 # make clean
8 # make menuconfig
9 # make -j10

```

5. 准备根文件系统

1. 我在 `/home/username/` 下建立了 `myfile` 文件夹，在这个文件夹下：

```

1 dd if=/dev/zero of=busyboxinitrd4M.img bs=4096 count=1024
2 mkfs.ext3 busyboxinitrd4M.img
3 mkdir rootfs
4 sudo mount -o loop busyboxinitrd4M.img rootfs/

```

2. 回到Busybox目录：输入如下命令：

```

1 make CONFIG_PREFIX= /home/ruizhao/myfile/rootfs/ install

```

3. 再次回到 `myfile` 目录下：

```

1 umount rootfs

```

6. 开始调试

```

1 ~/linux$ qemu-system-x86_64 -kernel ./linux-4.15.14/arch/x86_64/boot/bzImage
  ../myfile/busyboxinitrd4M.img -append "root=/dev/ram init=/bin/ash" -append nokaslr -s -S

```

下面是错误的历史：

第一次安装的QEMU版本旧的时候，就是这里出了问题，出现的界面如下

```

ruizhao@ruizhao-virtual-machine: ~/linux
ruizhao@ruizhao-virtual-machine:~$ cd linux/
ruizhao@ruizhao-virtual-machine:~/linux$ qemu-system-x86_64 -kernel ./linux-4.15
.14/arch/x86_64/boot/bzImage ../myfile/busyboxinitrd4M.img -append "root=/dev/ra
m init=/bin/ash" -append nokaslr -s -S
WARNING: Image format was not specified for '../myfile/busyboxinitrd4M.img' and
probing guessed raw.
    Automatically detecting the format is dangerous for raw images, write o
perations on block 0 will be restricted.
    Specify the 'raw' format explicitly to remove the restrictions.
warning: TCG doesn't support requested feature: CPUID.01H:ECX.vmx [bit 5]
main-loop: WARNING: I/O thread spun for 1000 iterations

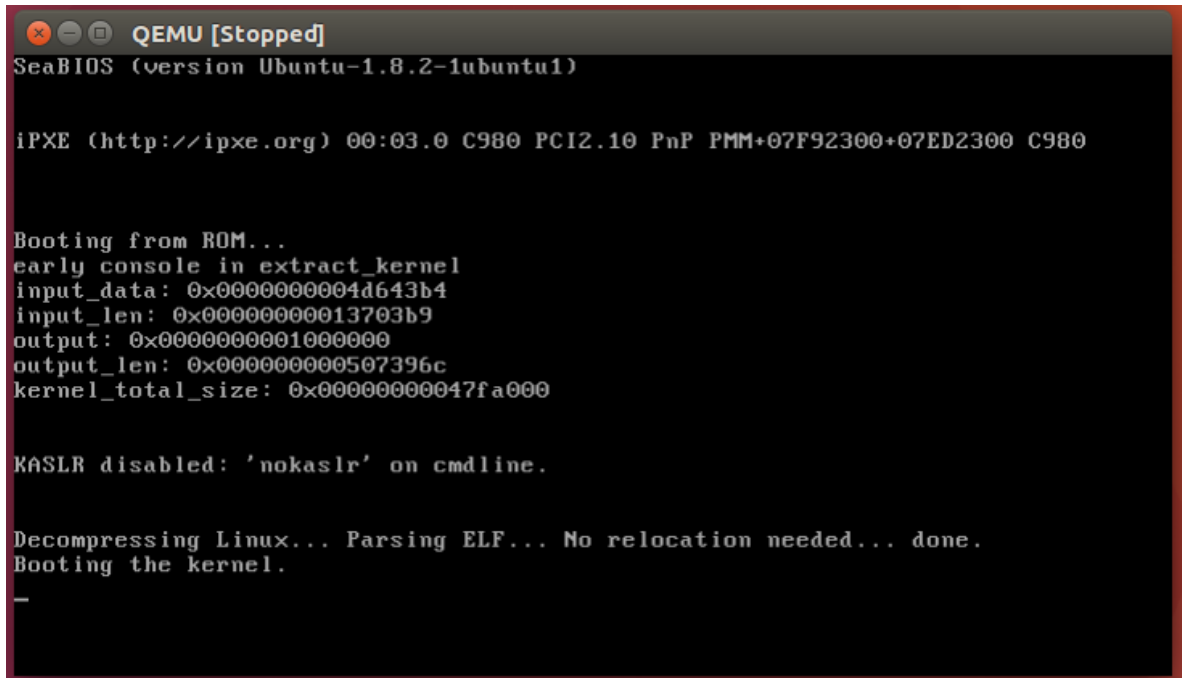
```

还出现了QEMU的界面，不过没有内容，此时一切正常，我再开一个terminal来GDB，

过程如下：

```
1 ~$ cd linux/  
2 ~/linux$ gdb  
3 (gdb) file linux-4.15.14/vmlinux  
4 (gdb) target remote:1234  
5 (gdb) break start_kernel  
6 (gdb) c
```

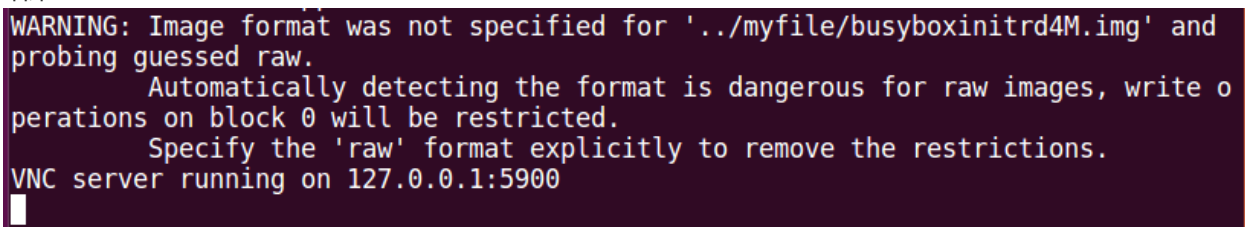
此时QEMU就显示了新内容，停在了start_kernel的部分。



```
QEMU [Stopped]  
SeaBIOS (version Ubuntu-1.8.2-1ubuntu1)  
  
iPXE (http://ipxe.org) 00:03.0 C980 PCI2.10 PnP PMM+07F92300+07ED2300 C980  
  
Booting from ROM...  
early console in extract_kernel  
input_data: 0x0000000004d643b4  
input_len: 0x00000000013703b9  
output: 0x0000000001000000  
output_len: 0x000000000507396c  
kernel_total_size: 0x00000000047fa000  
  
KASLR disabled: 'nokaslr' on cmdline.  
  
Decompressing Linux... Parsing ELF... No relocation needed... done.  
Booting the kernel.  
-
```

但是终端不可以用list来显示附近的代码，而且再次continue的话会显示该线程正在被使用，Google后有人说是GDB版本太久的bug，于是我更新了GDB，结果问题依然存在，后来同学也有这种问题，助教后来说是可以换Ubuntu17.10或者更新QEMU试试，我装了 Ubuntu17.10后又在前面的步骤遇到了新的问题，就回到Ubuntu16.04更新QEMU了，结果调试的问题小时了！看来真的是QEMU版本太旧的锅

结果：



```
WARNING: Image format was not specified for '../myfile/busyboxinitrd4M.img' and  
probing guessed raw.  
Automatically detecting the format is dangerous for raw images, write o  
perations on block 0 will be restricted.  
Specify the 'raw' format explicitly to remove the restrictions.  
VNC server running on 127.0.0.1:5900  
█
```

再打开一个terminal：使用GDB来调试：

```
1 $ gdb -tui
```

得到界面如图：

```
ruizhao@ruizhao-virtual-machine: ~  
[ No Source Available ]  
None No process In: L?? PC: ??  
Copyright (C) 2018 Free Software Foundation  
License GPLv3+: GNU GPL version 3 or later  
This is free software: you are free to chan  
There is NO WARRANTY, to the extent permitt  
and "show warranty" for details.  
This GDB was configured as "x86_64-pc-linux  
---Type <return> to continue, or q <return> to quit---
```

输入

```
1 return  
2 (gdb) file linux/linux-4.15.14/vmlinux  
3 (gdb) target remote:1234
```

```
remote Thread 1 In: cpu hw events L?? PC: 0xffff0
This GDB was configured as "x86_64-pc-linux
---Type <return> to continue, or q <return> to quit---return
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
(gdb) file linux/linux-4.15.14/vmlinux
Reading symbols from linux/linux-4.15.14/vmlinux...done.
(gdb) target remote:1234
Remote debugging using :1234
warning: Can not parse XML target description; XML support was disabled at compile time

0x00000000000000ffff in cpu_hw_events ()
(gdb) █
```

关键事件

1. 设置断点: `start_kernel`, 并开始运行

```
1 (gdb) break start_kernel
2 (gdb) c
```

```
ruizhao@ruizhao-virtual-machine: ~
init/main.c
506     vmalloc_init();
507     ioremap_huge_init();
508     /* Should be run before the first non-init thread is created */
509     init_espfix_bsp();
510     /* Should be run after espfix64 is set up. */
511     pti_init();
512 }
513
514 asmlinkage __visible void __init start_kernel(void)
B+> 515 {
516     char *command_line;
517     char *after_dashes;
518
519     set_task_stack_end_magic(&init_task);
520     smp_setup_processor_id();
521     debug_objects_early_init();
522
523     cgroup_init_early();
524
525     local_irq_disable();
526     early_boot_irqs_disabled = true;

remote Thread 1 In: start_kernel L515 PC: 0xffffffff854c4086
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
(gdb) file linux/linux-4.15.14/vmlinux
Reading symbols from linux/linux-4.15.14/vmlinux...done.
(gdb) target remote:1234
Remote debugging using :1234
warning: Can not parse XML target description; XML support was disabled at compile time

0x000000000000ffff0 in cpu_hw_events ()
(gdb) break start_kernel
Breakpoint 1 at 0xffffffff854c4086: file init/main.c, line 515.
(gdb) c
Continuing.

Breakpoint 1, start_kernel () at init/main.c:515
(gdb) □
```

2. 在 `start_kernel` 的初始之初你可以看到这两个变量:

```
1 char *command_line;
2 char *after_dashes;
```

第一个变量表示内核命令行的全局指针, 第二个变量将包含 `parse_args` 函数通过输入字符串中的参数 'name=value', 寻找特定的关键字和调用正确的处理程序。

3. 下一个函数是 `set_task_stack_end_magic`, 参数为 `init_task`。 `init_task` 代表初始化进程(任务)数据结构:

```
1 struct task_struct init_task = INIT_TASK(init_task);
```

`task_struct` 存储了进程的所有相关信息。

4. `set_task_stack_end_magic` 初始化完毕后的下一个函数是 `smp_setup_processor_id`。此函数在 `x86_64` 架构上是空函数:

```

1 void __init __weak smp_setup_processor_id(void)
2 {
3 }

```

在此架构上没有实现此函数

5. 然后运行到 `boot_cpu_init`，这里是激活第一个CPU事件

```

init/main.c
514 asmlinkage __visible void __init start_kernel(void)
515 {
516     char *command_line;
517     char *after_dashes;
518
519     set_task_stack_end_magic(&init_task);
520     smp_setup_processor_id();
521     debug_objects_early_init();
522
523     cgroup_init_early();
524
525     local_irq_disable();
526     early_boot_irqs_disabled = true;
527
528     /*
529      * Interrupts are still disabled. Do necessary setups, then
530      * enable them.
531      */
> 532     boot_cpu_init();
533     page_address_init();
534     pr_notice("%s", linux_banner);

```

remote Thread 1 In: start kernel L532 PC: 0xffffffff854c40c9
0x0000000000000000 in cpu_hw_events ()
(gdb) b start_kernel
Breakpoint 1 at 0xffffffff854c4086: file init/main.c, line 515.
(gdb) c
Continuing.

Breakpoint 1, start_kernel () at init/main.c:515
(gdb) b boot_cpu_init
Breakpoint 2 at 0xffffffff854fb3ac: file kernel/cpu.c, line 2008.
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n

进入这个函数观察：

```

kernel/cpu.c
2007 void __init boot_cpu_init(void)
B+> 2008 {
2009     int cpu = smp_processor_id();
2010
2011     /* Mark the boot cpu "present", "online" etc for SMP and UP case */
2012     set_cpu_online(cpu, true);
2013     set_cpu_active(cpu, true);
2014     set_cpu_present(cpu, true);
2015     set_cpu_possible(cpu, true);
2016
2017     #ifdef CONFIG_SMP
2018         __boot_cpu_id = cpu;
2019     #endif
2020 }
2021
2022 /*
2023  * Must be called _AFTER_ setting up the per_cpu areas
2024  */
2025 void __init boot_cpu_state_init(void)
2026 {
2027     per_cpu_ptr(&cpuhp_state, smp_processor_id())->state = CPUHP_ONLIN
remote Thread 1 In: boot_cpu_init L2008 PC: 0xffffffff854fb3ac
(gdb) c
Continuing.

Breakpoint 1, start_kernel () at init/main.c:515
(gdb) b boot_cpu_init
Breakpoint 2 at 0xffffffff854fb3ac: file kernel/cpu.c, line 2008.
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n
(gdb) n

Breakpoint 2, boot_cpu_init () at kernel/cpu.c:2008
(gdb)

```

首先我们需要获取当前处理器的ID通过下面函数：

```
1 int cpu = smp_processor_id();
```

现在是0.

6. Linux 内核的第一条打印信息：


```
init/main.c
525     local_irq_disable();
526     early_boot_irqs_disabled = true;
527
528     /*
529     * Interrupts are still disabled. Do necessary setups, then
530     * enable them.
531     */
532     boot_cpu_init();
533     page_address_init();
> 534     pr_notice("%s", linux_banner);
535     setup_arch(&command_line);
536     /*
537     * Set up the the initial canary and entropy after arch
538     * and after adding latent and command line entropy.
539     */
540     add_latent_entropy();
541     add_device_randomness(command_line, strlen(command_line));
542     boot_init_stack_canary();
543     mm_init_cpumask(&init_mm);
544     setup_command_line(command_line);
545     setup_nr_cpu_ids();

remote Thread 1 In: start_kernel                                L534  PC: 0xffffffff854c40ce
1823     asmlinkage __visible int printk(const char *fmt, ...)
1824     {
1825         va_list args;
1826         int r;
1827
1828         va_start(args, fmt);
1829         r = vprintk_func(fmt, args);
1830         va_end(args);
1831
1832         return r;
1833     }
```

调用了pr_notice函数。

```
1 #define pr_notice(fmt, ...) \
2     printk(KERN_NOTICE pr_fmt(fmt), ##__VA_ARGS__)
```

pr_notice其实是printk的扩展，这里我们使用它打印了Linux的banner。

```
1 pr_notice("%s", linux_banner);
```

打印的是内核的版本号以及编译环境信息。

7. 依赖于体系结构的初始化部分

```

init/main.c
527
528      /*
529       * Interrupts are still disabled. Do necessary setups, then
530       * enable them.
531       */
532      boot_cpu_init();
533      page_address_init();
534      pr_notice("%s", linux_banner);
535      setup_arch(&command_line);
536      /*
537       * Set up the the initial canary and entropy after arch
538       * and after adding latent and command line entropy.
539       */
540      add_latent_entropy();

```

8. rest_init()

这是 start_kernel 的最后一个函数

```

@ruizhao-virtual-machine: ~/linux
init/main.c
698      proc_root_init();
699      nsfs_init();
700      cpuset_init();
701      cgroup_init();
702      taskstats_init_early();
703      delayacct_init();
704
705      check_bugs();
706
707      acpi_subsystem_init();
708      arch_post_acpi_subsys_init();
709      sfi_init_late();
710
711      if (efi_enabled(EFI_RUNTIME_SERVICES)) {
712          efi_free_boot_services();
713      }
714
715      /* Do the rest non-__init'ed, we're now alive */
> 716      rest_init();
717  }
718

```

```

@ruizhao-virtual-machine: ~/linux
init/main.c
391      static noinline void __ref rest_init(void)
> 392  {
393          struct task_struct *tsk;
394          int pid;
395
396          rcu_scheduler_starting();
397          /*
398           * We need to spawn init first so that it obtains pid 1, however
399           * the init task will end up wanting to create kthreads, which, if
400           * we schedule it before we create kthreadd, will OOPS.
401           */
402          pid = kernel_thread(kernel_init, NULL, CLONE_FS);
403          /*
404           * Pin init on the boot CPU. Task migration is not properly working
405           * until sched_init_smp() has been run. It will set the allowed
406           * CPUs for init to the non isolated CPUs.
407           */
408          rcu_read_lock();
409          tsk = find_task_by_pid_ns(pid, &init_pid_ns);
410          set_cpus_allowed_ptr(tsk, cpumask_of(smp_processor_id()));
411          rcu_read_unlock();

```

remote Thread 1 In: rest_init L392 PC: 0xffffffff82a9a520

调用 `rest_init()` 函数进行最后的初始化工作,包括创建1号进程 (`init`),第一个内核线程等操作。最后,初始化结束。

实验总结

这次实验过程很艰辛,刚开始QEMU的版本问题和编译内核时没有选debug info等导致了前期花费了大量时间来搭建调试环境。通过这次实验,对Linux和Linux内核不再陌生了,刚开始做实验的时候都不知道从何看起,熟悉了Linux的一些命令,学习了一些GDB的调试命令,虽然很艰辛,但是感觉很有收获。